

Lab 6: Copy-on-Write Fork in xv6

Report

Tejesh Raut
140050008

Gangesh Gudmalwar
140050058

October 2, 2016

To run the program: Clone the git repo of xv6 and then patch our changes into it by entering the following commands:

```
git clone https://github.com/guilleiguaran/xv6.git
patch -p0 -i file.patch
cd xv6
make
make qemu
```

Then to run testcase enter following command in qemu:

```
testcow
```

Project Report:

1 Adding getNumFreePages system call

Add the variable `numFreePages` to implement the given system call in `kalloc.c` :

In `kalloc.c` add the entry in struct `kmem`:

```
int numFreePages;
```

Adding given system calls

Add following code in `defs.h`:

```
int getNumFreePages(void);
```

Add following code in `sysproc.c`:

```
int
sys_getNumFreePages(void)
{
```

```
    return getNumFreePages();  
}
```

Add the following declaration along with other system calls in syscall.c:

```
extern int sys_getNumFreePages(void);
```

Add following fields in the same file like other system calls in syscall.c:

```
[SYS_getNumFreePages] sys_getNumFreePages,
```

Add the following lines in usys.S:

```
SYSCALL(getNumFreePages)
```

Add following lines in syscall.h:

```
#define SYS_getNumFreePages 22
```

Add the function body of getNumFreePages in kalloc.c

```
// Returns the number of free pages.  
int  
getNumFreePages(void)  
{  
    if(kmem.use_lock)  
        acquire(&kmem.lock);  
    int r = kmem.numFreePages;  
    if(kmem.use_lock)  
        release(&kmem.lock);  
    return (r);  
}
```

Initialize numFreePages in kinit1

```
kmem.numFreePages = 0;
```

In kfree

```
kmem.numFreePages = kmem.numFreePages + 1; // A new node is added to  
    freelist so increase the number of free pages
```

In kalloc

```
kmem.numFreePages = kmem.numFreePages - 1 ; // A node is popped out  
    from the freelist so decrease the number of free pages.
```

2 Reinstalling of page table

Whenever the flags are changed in copyuvm function the page table must be reinstalled using:

```
lcr3(v2p(pgdir)); // reinstall the page table
```

3 Keeping track of the reference count of pages

In vm.c add the declaration of array and lock:

```
struct spinlock lock;  
char pg_refcount[PHYSTOP >> PGSHIFT]; // array to store refcount
```

In inituvm function:

```
acquire(&lock);  
pg_refcount[v2p(mem) >> PGSHIFT] = pg_refcount[v2p(mem) >> PGSHIFT] +  
    1 ;  
release(&lock);
```

In allocuvm function:

```
acquire(&lock);  
pg_refcount[v2p(mem) >> PGSHIFT] = pg_refcount[v2p(mem) >> PGSHIFT] +  
    1 ;  
release(&lock);
```

In deallocuvm free the page only when no other page table is pointing it:

```
acquire(&lock);  
if(--pg_refcount[pa >> PGSHIFT] == 0) // if no other page table is  
    pointing to this page remove it  
{  
    char *v = p2v(pa);  
    kfree(v);  
}  
release(&lock);
```

In copyuvm function when a process is forked the refcount of that permanent address should be incremented:

```
acquire(&lock);  
pg_refcount[pa >> PGSHIFT] = pg_refcount[pa >> PGSHIFT] + 1; //  
    increase reference count of that permanent page.  
release(&lock);
```

4 Change of copyuvm function

Make the pagetable unwritable and then assign the same permanent addresses to the new page table

```
pde_t*
copyuvm(pde_t *pgdir, uint sz)
{
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    //char *mem; //No need to allocate new memory

    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyuvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyuvm: page not present");
        *pte &= ~PTE_W; // make this page table unwritable
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        // No need of page allocation
        if(mappages(d, (void*)i, PGSIZE, pa, flags) < 0) // map the
            child's page table to same permanent addresses
            goto bad;
        acquire(&lock);
        pg_refcount[pa >> PGSHIFT] = pg_refcount[pa >> PGSHIFT] + 1; //
            increase reference count of that permanent page.
        release(&lock);
    }
    lcr3(v2p(pgdir)); // reinstall the page table
    return d;

bad:
    freevm(d);
    lcr3(v2p(pgdir)); // reinstall the page table
    return 0;
}
```

5 Adding trap handler to handle pagefaults

To show error on page fault In defs.h

```
void    pagefault(uint err_code);
```

In trap.c:

```
case T_PGFLT:
    pagefault(tf->err);
    break;
```

In vm.c

```
void pagefault(uint err_code)
{
    cprintf("Pagefault occured");
    return;
}
```

6 Adding trap handling function to make copy of user memory

In vm.c

```
void pagefault(uint err_code)
{
    uint va = rcr2();
    uint pa;
    pte_t *pte;
    char *mem;
    if(va >= KERNBASE)
    {
        cprintf("pid %d %s: Illegal memory access on CPU %d due to
                virtual address 0x%x is mapped to kernel code. So killing
                the process\n", proc->pid, proc->name, cpu->id, va);
        proc->killed = 1;
        return;
    }
    if((pte = walkpgdir(proc->pgdir, (void*)va, 0))==0)
    {
        cprintf("pid %d %s: Illegal memory access on CPU %d due to
                virtual address 0x%x is mapped to NULL pte. So killing the
                process\n", proc->pid, proc->name, cpu->id, va);
        proc->killed = 1;
        return;
    }
}
```

```

}
if(!(*pte & PTE_P))
{
    cprintf("pid %d %s: Illegal memory access on CPU %d due to
        virtual address 0x%x is mapped to pte which is not present.
        So killing the process\n", proc->pid, proc->name, cpu->id,
        va);
    proc->killed = 1;
    return;
}
if(!(*pte & PTE_U))
{
    cprintf("pid %d %s: Illegal memory access on CPU %d due to
        virtual address 0x%x is mapped to pte which is not
        accessible to user. So killing the process\n", proc->pid,
        proc->name, cpu->id, va);
    proc->killed = 1;
    return;
}
if(*pte & PTE_W)
{
    panic("Unknown page fault due to a writable pte");
}
else
{
    pa = PTE_ADDR(*pte);
    acquire(&lock);
    if(pg_refcount[pa >> PGSHIFT] == 1)
    {
        release(&lock);
        *pte |= PTE_W;
    }
    else
    {
        if(pg_refcount[pa >> PGSHIFT] > 1)
        {
            release(&lock);
            if((mem = kalloc()) == 0)
            {
                cprintf("pid %d %s: Pagefault due to out of memory",
                    proc->pid, proc->name);
                proc->killed = 1;
                return;
            }
            memmove(mem, (char*)p2v(pa), PGSIZE);
            acquire(&lock);
            pg_refcount[pa >> PGSHIFT] = pg_refcount[pa >> PGSHIFT] -
                1;
            pg_refcount[v2p(mem) >> PGSHIFT] = pg_refcount[v2p(mem) >>
                PGSHIFT] + 1;

```

```

        release(&lock);
        *pte = v2p(mem) | PTE_P | PTE_W | PTE_U;
    }
    else
    {
        release(&lock);
        panic("Pagefault due to wrong ref count");
    }
}
lcr3(v2p(proc->pgdir));
}
}

```

7 Test case testcow

```

#include "types.h"
#include "stat.h"
#include "user.h"

int i = 3;
int main(void)
{
    int pid;
    pid = fork();
    if(pid == 0)
    {
        printf(1, "Number of free pages in child 1 before changing\n", getNumFreePages());
        i = 4;
        printf(1, "Number of free pages in child 1 after changing\n", getNumFreePages());
    }
    else
    {
        wait();
        pid = fork();
        if(pid == 0)
        {
            printf(1, "Number of free pages in child 2 before changing\n", getNumFreePages());
            i = 4;
            printf(1, "Number of free pages in child 2 after changing\n", getNumFreePages());
        }
        else
        {

```

```
        printf(1,"Number of free pages in parent are: %d\n",
               getNumFreePages());
        wait();
    }
}
exit();
}
```

Output of above test case is :

```
$ testcow
Number of free pages in child 1 before changing variable are: 56710
Number of free pages in child 1 after changing variable due to copy
are: 56709
Number of free pages in parent are: 56710
Number of free pages in child 2 before changing variable are: 56710
Number of free pages in child 2 after changing variable due to copy
are: 56709
```
